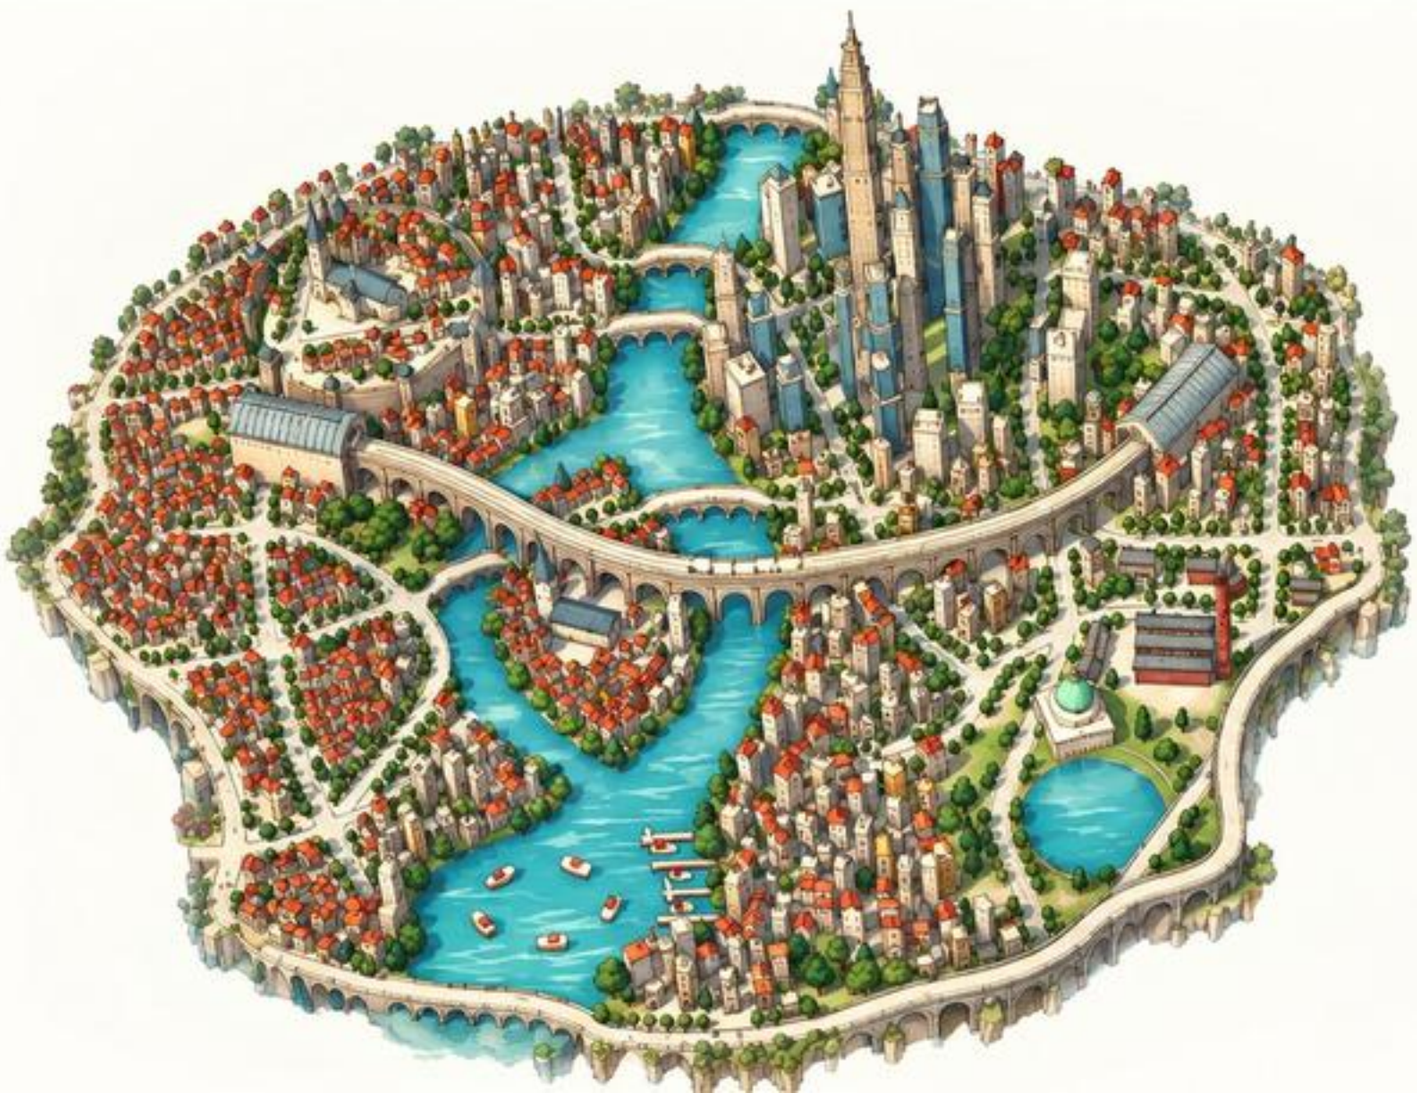


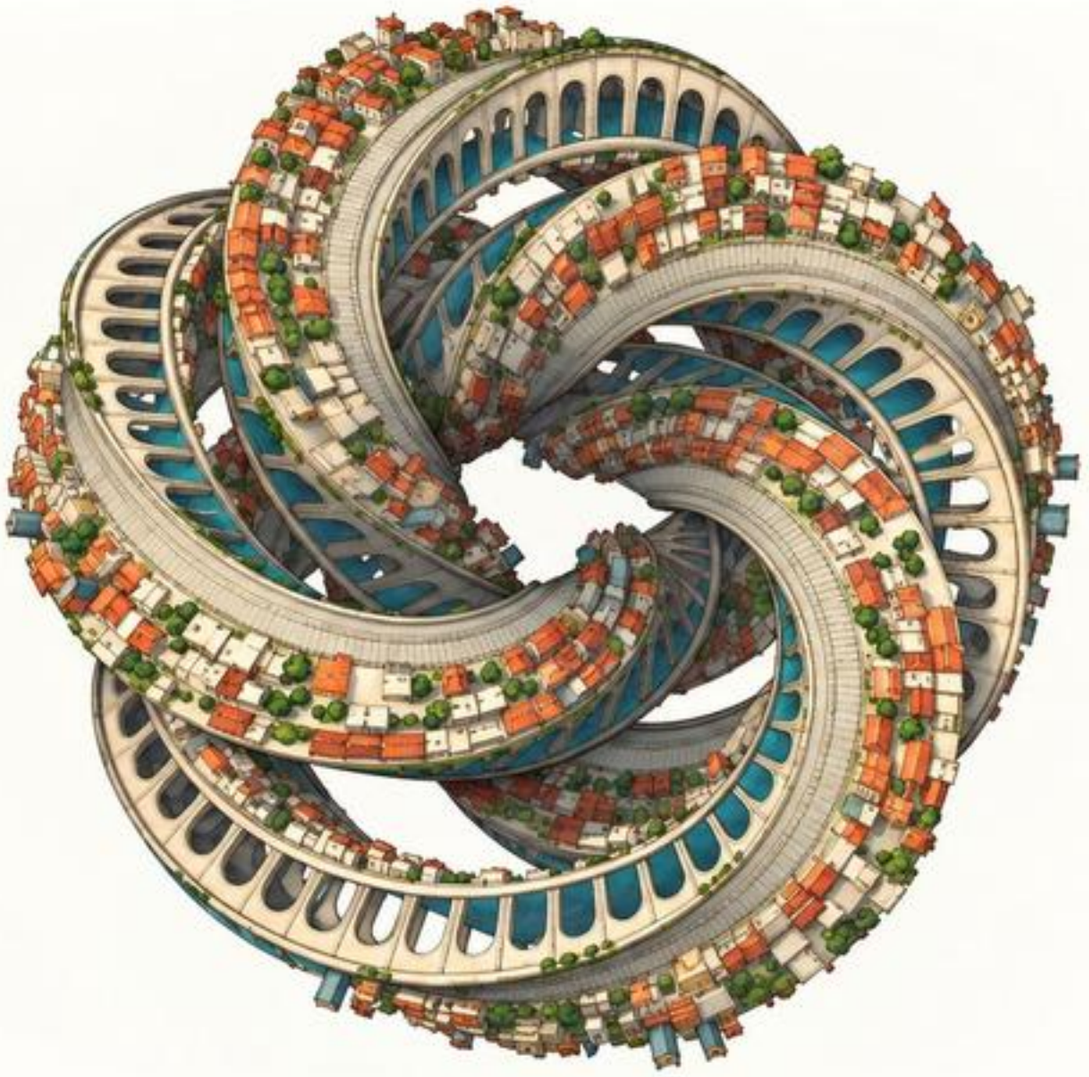
arxdsilva
@arxdsilva

Goodfire says 90% of neurons in a 32B model have one main job.

They got there by fixing ordinary problems the field had blamed on superposition.

Nick Cammarata calls it a working hypothesis from a preliminary audit.







arxdsilva
@arxdsilva

A model you can read is a model you can debug.

One researcher in the discussion isn't ready to drop superposition.

Which step would you skip? Save this.

- Longer contexts: short ones made long-context neurons look mixed
- Name the shared pattern: the cat-and-car neuron detected a shape
- Residuals and LayerNorm kept deep layers from overflowing